

Instructions

In this assignment you will build recommender systems to make predictions related to book reviews from *Goodreads*.

Submissions will take the form of prediction files uploaded to gradescope, where their test set performance will be evaluated on a leaderboard. **Most of your grade will be determined by ‘absolute’ cutoffs; the leaderboard ranking will only determine enough of your assignment grade to make the assignment FUN.**

The assignment is due at the beginning of week 8, though **the exact deadline is as posted on gradescope**. Make sure you upload solutions to the leaderboard regularly in order to make continual progress.

You should submit two files:

- `writeup.txt` a brief, plain-text description of your solutions to each task; please prepare this adequately in advance of the submission deadline; this is only intended to help us follow your code and does not need to be detailed.
- `assignment1.py` A python file containing working code for your solutions. *The autograder will not execute your code*; this file is required so that we can assign partial grades in the event of incorrect solutions, check for plagiarism, etc. Your solution should clearly document which sections correspond to each task. We may occasionally run code to confirm that your outputs match submitted answers, so please ensure that your code generates the submitted answers.

Along with three files corresponding to your predictions:

`predictions Read.csv`, `predictions Category.csv`, `predictions Rating.csv`

These files should contain your predictions for each (test) instance. **The provided baseline code demonstrates how to generate valid output files.**

This assignment should be completed individually. To begin, download the files for this assignment from:

https://drive.google.com/file/d/1XeWAovw-LKbToPL_STHoeCgDEYvu9yhX/view?usp=sharing

Files

- `train_Interactions.csv.gz` 200,000 ratings to be used for training. This data should be used for the ‘read prediction’ and ‘rating prediction’ tasks. It is not necessary to use all ratings for training, for example if doing so proves too computationally intensive.
 - `userID` The ID of the user. This is a hashed user identifier from Goodreads.
 - `bookID` The ID of the book. This is a hashed book identifier from Goodreads.
 - `rating` The star rating of the user’s review.

- `train_Category.json.gz` Training data for the category prediction task. This file is json formatted, and contains the following fields:
 - `n votes` The number of ‘likes’ this review has received.
 - `review id` A hashed identifier for this review.
 - `user id` A hashed identifier for the user.
 - `review text` Text of the review.
 - `rating` Rating of the book.
 - `genreID` A numeric label associated with the genre.
 - `genre` A string version of the genre.
- `test_Category.json.gz` Test data associated with the category prediction task. This data has the same format as above, with the ‘genre’ and ‘genreID’ labels hidden.
- `pairs_Read.csv` Pairs on which you are to predict whether a book was read (both classes).
- `pairs_Category.csv` Pairs (userID and reviewID) on which you are to predict the category of a book.
- `pairs_Rating.csv` Pairs (userIDs and bookIDs) on which you are to predict ratings.
- `baselines.py` A simple baseline for each task, described below.

Please do not try to collect these reviews from Goodreads, or to reverse-engineer the hashing function I used to anonymize the data. Doing so will not be easier than successfully completing the assignment. **We may execute code for any solution suspected of violating the competition rules to confirm that it generates valid output;** code will be run through a plagiarism detector. **We may assign a zero grade if your code does not appear to correspond to your submitted solution.**

Tasks

You are expected to complete the following tasks:

- **Read prediction:** Predict given a (user,book) pair from ‘pairs Read.csv’ whether the user would read the book (0 or 1). Accuracy will be measured in terms of the categorization accuracy (fraction of correct predictions). The test set has been constructed such that exactly 50% of the pairs correspond to read books and the other 50% do not.
- **Category prediction:** Predict the category of a book from a review. Five categories are used for this task, which can be seen in the baseline program, namely Children’s, Comics/Graphic Novels, Fantasy, Mystery/Thriller, and Romance. Performance will be measured in terms of the fraction of correct classifications (this is admittedly not really a recommender systems task, and is more related to classification / feature engineering).
- **Rating prediction:** Predict people’s star ratings as accurately as possible, for those (user,item) pairs in ‘pairs Rating.txt’. Accuracy will be measured in terms of the mean-squared error (MSE).

A competition page has been set up on gradescope to keep track of your results compared to those of other members of the class. The leaderboard will show your results on half of the test data, but your ultimate score will depend on your predictions across the whole dataset.

Grading and Evaluation

This assignment is worth 27% of your grade. You will be graded on the following aspects. Each of the three tasks is worth 8 marks, plus 3 marks for the written report. Your grade for each task will be based on:

- Your ability to obtain a solution which outperforms the leaderboard baselines on the *unseen portion* of the test data (4 marks for each task). Obtaining full marks requires a solution which is substantially better than baseline performance.
- Your ranking for each of the tasks compared to other students in the class (2 marks for each task).
- Obtain a solution which outperforms the baselines on the *seen portion* of the test data (i.e., the leaderboard). This is a sort of consolation prize in case you overfit to the leaderboard (2 marks for each task).

Finally, your written report should describe the approaches you took to each of the tasks. Your report should contain a few sentences describing your solution to each task. You will generally get full marks for the report automatically unless some issue is suspected with your solution.

To obtain good performance, you should not need to invent new approaches (though you are more than welcome to!) but rather you will be graded based on your decision to apply reasonable approaches to each of the given tasks.

Baselines

Simple baselines have been provided for each of the tasks. These are included in ‘`baselines.py`’ among the files above. They are mostly intended to demonstrate how the data is processed and prepared for submission to gradescope. These baselines operate as follows:

- **Read prediction:** Find the most popular books that account for 50% of interactions in the training data. Return “1” whenever such a book is seen at test time, “0” otherwise.
- **Category prediction:** Look for a few likely words that may appear in reviews of each category (e.g. if the word “fantasy” appears, classify as Fantasy).
- **Rating prediction:** Return the global average rating, or the user’s average if we have seen them before in the training data.

Running ‘`baselines.py`’ produces files containing predicted outputs (these outputs can be uploaded to gradescope). Your submission files should have the same format.